

Lab 14

Reinforcement Learning

Datalab

Department of Computer Science,
National Tsing Hua University, Taiwan

Outline

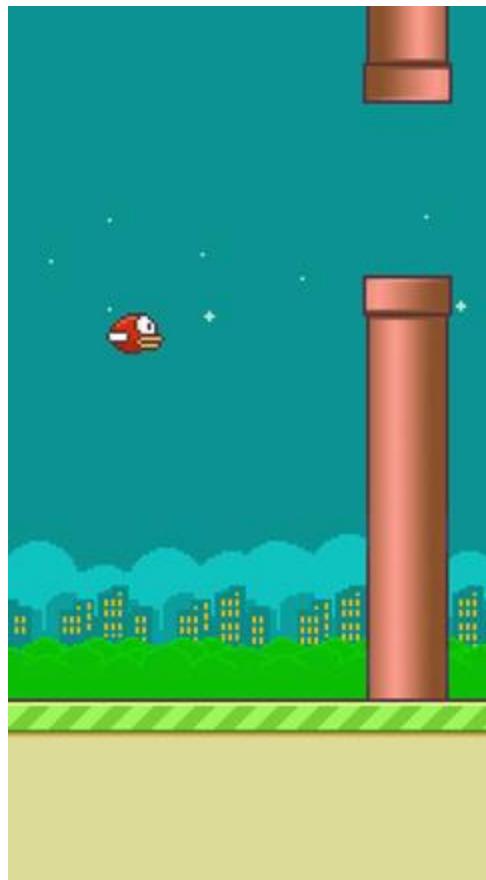
- Homework
- Markov Decision Process (MDP)
 - Value Iteration
 - Policy Iteration
- Q-Learning & SARSA

Outline

- Homework
 - MDP(value iteration & policy iteration)
 - Q-Learning & SARSA

Homework

- Train an agent to play Flappy Bird game(SARSA)



Install PLE and Pygame

- Clone the repo

```
$ git clone https://github.com/ntasfi/PyGame-Learning-Environment
Cloning into 'PyGame-Learning-Environment'...
remote: Enumerating objects: 1118, done.
remote: Total 1118 (delta 0), reused 0 (delta 0), pack-reused 1118
Receiving objects: 100% (1118/1118), 8.06 MiB | 800.00 KiB/s, done.
Resolving deltas: 100% (592/592), done.
```

- Install PLE(in the PyGame-Learning-Environment folder)

- cd PyGame-Learning-Environment
 - pip install –e .

```
$ pip install -e .
Obtaining file:///E:/DL/RL/PyGame-Learning-Environment
Requirement already satisfied: numpy in c:\users\vincent\anaconda3\lib\site-packages (from ple==0.0.1) (1.16.4)
Requirement already satisfied: Pillow in c:\users\vincent\anaconda3\lib\site-packages (from ple==0.0.1) (6.1.0)
Installing collected packages: ple
  Found existing installation: ple 0.0.1
    Uninstalling ple-0.0.1:
      Successfully uninstalled ple-0.0.1
    Running setup.py develop for ple
Successfully installed ple
```

- pip install pygame

Homework

- What you should do:
 - Change the update rule from Q-learning to **SARSA (with the same episodes)**.
 - Give a brief report to discuss the result (compare Q-learning with SARSA based on the game result).
- Remind
 - Only need **CPU** resources.
 - It will take you more than **13** hours to train, please reserve enough time.

Homework

- Precautions:
 - If you encounter this problem, just stop.
 - It means your bird plays well and the recorded frames is too long to save.

```
~\Anaconda3\lib\site-packages\moviepy\video\io\html_tools.py in html_embed(clip, filetype, maxduration, rd_kwargs, center, **html_k
wargs)
 105
 106     return html_embed(filename, maxduration=maxduration, rd_kwargs=rd_kwargs,
--> 107             center=center, **html_kwargs)
 108
 109     filename = clip

~\Anaconda3\lib\site-packages\moviepy\video\io\html_tools.py in html_embed(clip, filetype, maxduration, rd_kwargs, center, **html_k
wargs)
 140     if duration > maxduration:
 141         raise ValueError("The duration of video %s (%.1f) exceeds the 'maxduration' "%(filename, duration)+
--> 142             "attribute. You can increase 'maxduration', by passing 'maxduration' parameter"
 143                 "to ipython_display function."
 144             "But note that embedding large videos may take all the memory away !")

ValueError: The duration of video _temp_.mp4 (129.8) exceeds the 'maxduration' attribute. You can increase 'maxduration', by passing 'max
duration' parameter to ipython_display function. But note that embedding large videos may take all the memory away !
```

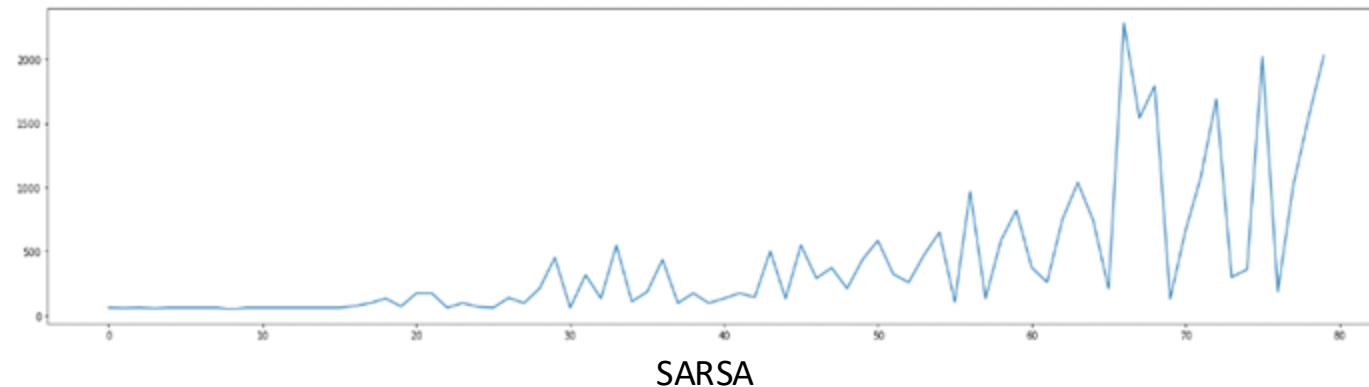
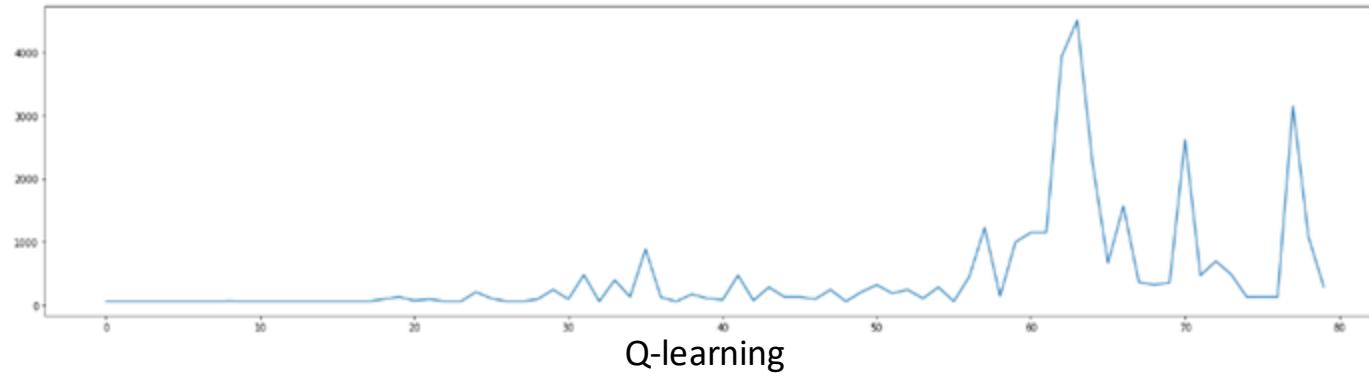
Homework

Requirements

- Write a brief report in the notebook
- Upload both `ipynb` and `mp4` to google drive
 - `Lab14_{student_id}.ipynb` (90%)
 - `Lab14_{student_id}.mp4` (10%)
- Share your drive's link via `eeclass`
 - Please make sure that TA can access your google drive!!!
(Or you will get 0 on this lab!)
- Deadline: 2025-12-03(Wed) 23:59

Homework

- Requirement (report):
 - You can compare life time or reward against training episodes.



Outline

- Homework
- Markov Decision Process (MDP)
 - Value Iteration
 - Policy Iteration
- Q-Learning & SARSA

Markov Decision Process (MDP)

- A MDP is defined by

S	A	P	R	γ	H
State space	Action space	Transition Probability	Reward	Discount Factor	Horizon

 小吃	資電	排球
綜二	台達	籃球
總圖	工三	西門

$S = \{ \text{小吃}, \text{資電}, \text{排球}, \text{綜二}, \text{台達}, \text{籃球}, \text{總圖}, \text{工三}, \text{西門} \}$

$A = \{ \text{上}, \text{下}, \text{左}, \text{右} \}$

$P = \text{no noise}$

$R(\text{工三}) = 0, R(\text{others}) = -1$

$r = 1$

$H = \inf$

We have a MDP model, then?

Goal - Find the Optimal Policy

- If the agent follow the optimal policy, it will get maximal total reward
- We can solve it via these two algorithms
 - Value Iteration
 - Policy Iteration

Outline

- Homework
- Markov Decision Process (MDP)
 - Value Iteration
 - Policy Iteration
- Q-Learning & SARSA

Value Iteration

Input: MDP $(\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty)$

Output: $\pi^*(s)$'s for all s 's

For each state s , initialize $V^*(s) \leftarrow 0$;

repeat

foreach s **do**

$| V^*(s) \leftarrow \max_a \sum_{s'} P(s'|s; a)[R(s, a, s') + \gamma V^*(s')];$
 end

until $V^*(s)$'s converge;

foreach s **do**

$| \pi^*(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a)[R(s, a, s') + \gamma V^*(s')];$

end



小吃	資電	排球
綜二	台達	籃球
總圖	工三	西門

$S = \{ \text{小吃}, \text{資電}, \text{排球}, \text{綜二}, \text{台達}, \text{籃球}, \text{總圖}, \text{工三}, \text{西門} \}$

$A = \{ \text{上}, \text{下}, \text{左}, \text{右} \}$

$P = \text{no noise}$

$R(\text{工三}) = 0, R(\text{others}) = -1$

$r = 1$

$H = \inf$

Input: MDP $(\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty)$ ←

Output: $\pi^*(s)$'s for all s 's

For each state s , initialize $V^*(s) \leftarrow 0$;
repeat

foreach s do
| $V^*(s) \leftarrow \max_a \sum_{s'} P(s'|s; a)[R(s, a, s') + \gamma V^*(s')]$;
| end

until $V^*(s)$'s converge;

foreach s do
| $\pi^*(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a)[R(s, a, s') + \gamma V^*(s')]$;
| end

After Initialization

	小吃 0	資電 0	排球 0
綜二 0	台達 0	籃球 0	
總圖 0	工三 0	西門 0	

Input: MDP $(\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty)$

Output: $\pi^*(s)$'s for all s 's

For each state s , initialize $V^*(s) \leftarrow 0$; 
repeat

foreach s do
| $V^*(s) \leftarrow \max_a \sum_{s'} P(s'|s; a)[R(s, a, s') + \gamma V^*(s')]$;
| end

until $V^*(s)$'s converge;

foreach s do
| $\pi^*(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a)[R(s, a, s') + \gamma V^*(s')]$;
| end

After Iteration 1

	小吃 -1	資電 -1	排球 -1
綜二 -1	台達 -1	籃球 -1	
總圖 -1	工三 0	西門 -1	

Input: MDP $(\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty)$

Output: $\pi^*(s)$'s for all s 's

For each state s , initialize $V^*(s) \leftarrow 0$;

repeat

foreach s do

$V^*(s) \leftarrow \max_a \sum_{s'} P(s'|s; a)[R(s, a, s') + \gamma V^*(s')]$;
end

until $V^*(s)$'s converge;

foreach s do

$\pi^*(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a)[R(s, a, s') + \gamma V^*(s')]$;
end

$$V(\text{台達}) = \text{Reward} + V(\text{資電}) = -1 + 0$$

$$V(\text{台達}) = \text{Reward} + V(\text{綜二}) = -1 + 0$$

$$V(\text{台達}) = \text{Reward} + V(\text{工三}) = -1 + 0$$

$$V(\text{台達}) = \text{Reward} + V(\text{籃球}) = -1 + 0$$

After Iteration 2



小吃 -2	資電 -2	排球 -2
綜二 -2	台達 -1	籃球 -2
總圖 -1	工三 0	西門 -1

Input: MDP $(\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty)$

Output: $\pi^*(s)$'s for all s 's

For each state s , initialize $V^*(s) \leftarrow 0$;

repeat

foreach s do

$| V^*(s) \leftarrow \max_a \sum_{s'} P(s'|s; a)[R(s, a, s') + \gamma V^*(s')]$;
| end

until $V^*(s)$'s converge;

foreach s do

$| \pi^*(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a)[R(s, a, s') + \gamma V^*(s')]$;
| end

$$V(\text{台達}) = \text{Reward} + V(\text{資電}) = -1 + -1$$

$$V(\text{台達}) = \text{Reward} + V(\text{綜二}) = -1 + -1$$

$$\max_a V(\text{台達}) = \text{Reward} + V(\text{工三}) = -1 + 0$$

$$V(\text{台達}) = \text{Reward} + V(\text{籃球}) = -1 + -1$$

After Iteration 3



小吃 -3	資電 -2	排球 -3
綜二 -2	台達 -1	籃球 -2
總圖 -1	工三 0	西門 -1

$$V(\text{小吃}) = V(\text{排球}) = -1 + -2 = -3$$

Input: MDP $(\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty)$

Output: $\pi^*(s)$'s for all s 's

For each state s , initialize $V^*(s) \leftarrow 0$;

repeat

foreach s do

$V^*(s) \leftarrow \max_a \sum_{s'} P(s'|s; a) [R(s, a, s') + \gamma V^*(s')]$;

end

until $V^*(s)$'s converge;

foreach s do

$\pi^*(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a) [R(s, a, s') + \gamma V^*(s')]$;

end

After Iteration 4

 小吃 -3	資電 -2	排球 -3
綜二 -2	台達 -1	籃球 -2
總圖 -1	工三 0 	西門 -1

Input: MDP $(\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty)$

Output: $\pi^*(s)$'s for all s 's

For each state s , initialize $V^*(s) \leftarrow 0$;

repeat

 foreach s do

$V^*(s) \leftarrow \max_a \sum_{s'} P(s'|s; a)[R(s, a, s') + \gamma V^*(s')]$;
 end

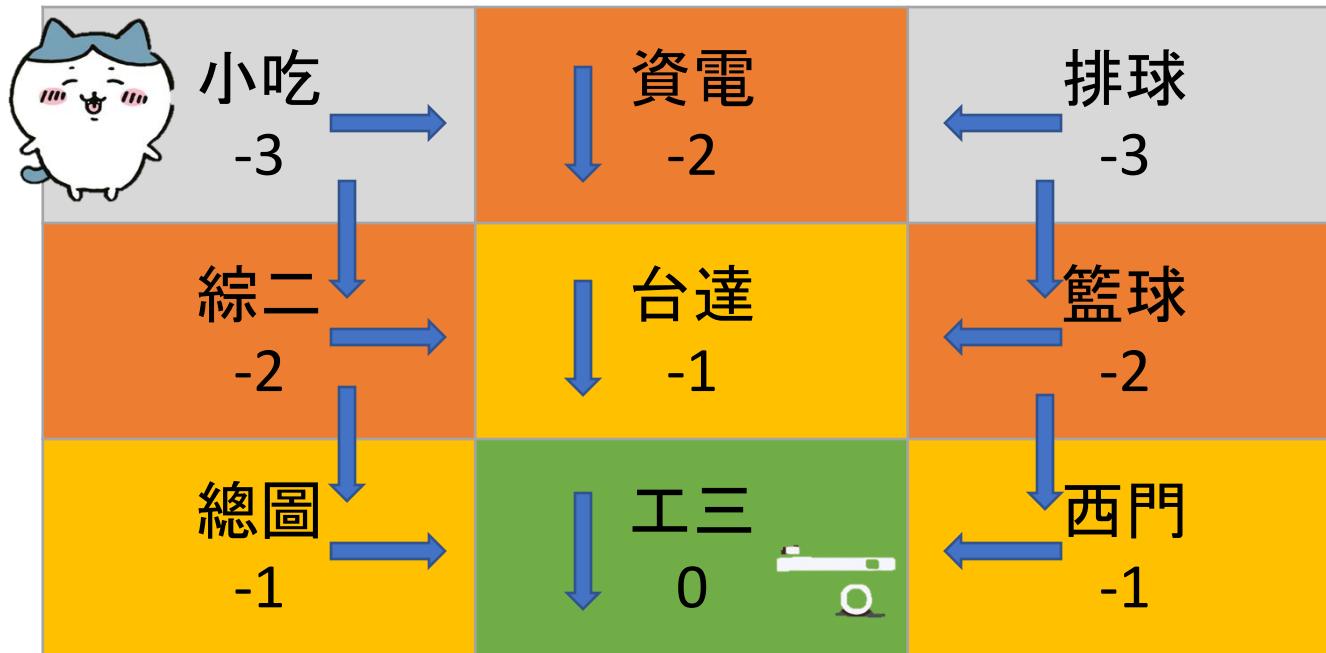
until $V^*(s)$'s converge; 

foreach s do

$\pi^*(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a)[R(s, a, s') + \gamma V^*(s')]$;
end

Iteration 4 = Iteration 3

Converge !



Input: MDP $(\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty)$

Output: $\pi^*(s)$'s for all s 's

For each state s , initialize $V^*(s) \leftarrow 0$;
repeat

```

foreach s do
    |  $V^*(s) \leftarrow \max_a \sum_{s'} P(s'|s; a)[R(s, a, s') + \gamma V^*(s')]$ ;
end

```

until $V^*(s)$'s converge;

foreach s do

```

    |  $\pi^*(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a)[R(s, a, s') + \gamma V^*(s')]$ ;
end

```

Now we have the optimal policy!



Outline

- Homework
- Markov Decision Process (MDP)
 - Value Iteration
 - Policy Iteration
- Q-Learning & SARSA

Policy Iteration

Input: MDP $(\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty)$

Output: $\pi(s)$'s for all s 's

For each state s , initialize $\pi(s)$ randomly;

repeat

 For each state s , initialize $V_\pi(s) \leftarrow 0$;

repeat

foreach s **do**

$| V_\pi(s) \leftarrow \sum_{s'} P(s'|s; \pi(s)) [R(s, \pi(s), s') + \gamma V_\pi(s')];$

end

until $V_\pi(s)$'s converge;

foreach s **do**

Policy evaluation

$| \pi(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a) [R(s, a, s') + \gamma V_\pi(s')];$

end

until $\pi(s)$'s converge;



小吃

資電

排球

綜二

台達

籃球

總圖

工三



西門

$S = \{ \text{小吃}, \text{資電}, \text{排球}, \text{綜二}, \text{台達}, \text{籃球}, \text{總圖}, \text{工三}, \text{西門} \}$

$A = \{ \text{上}, \text{下}, \text{左}, \text{右} \}$

$P = \text{no noise}$

$R(\text{工三}) = 0, R(\text{others}) = -1$

$r = 1$

$H = \inf$

Input: MDP ($\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty$)



Output: $\pi(s)$'s for all s 's

For each state s , initialize $\pi(s)$ randomly;
repeat

 For each state s , initialize $V_\pi(s) \leftarrow 0$;

 repeat

 foreach s do

$| V_\pi(s) \leftarrow \sum_{s'} P(s'|s; \pi(s)) [R(s, \pi(s), s') + \gamma V_\pi(s')];$
 end

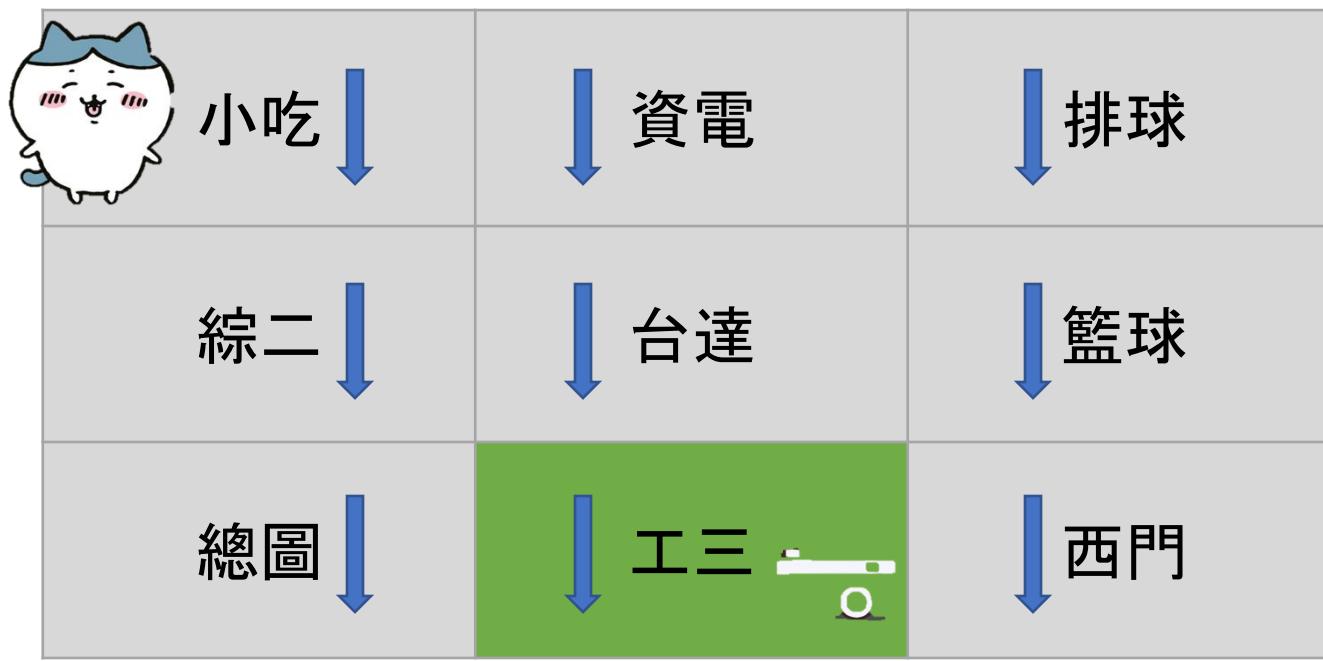
 until $V_\pi(s)$'s converge;

 foreach s do

 Policy evaluation

$| \pi(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a) [R(s, a, s') + \gamma V_\pi(s')];$
 end

 until $\pi(s)$'s converge;



Input: MDP ($\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty$)

Output: $\pi(s)$'s for all s's

For each state s , initialize $\pi(s)$ randomly;
repeat

 For each state s , initialize $V_\pi(s) \leftarrow 0$;

 repeat

 foreach s do

$| V_\pi(s) \leftarrow \sum_{s'} P(s'|s; \pi(s)) [R(s, \pi(s), s') + \gamma V_\pi(s')];$

 | end

 until $V_\pi(s)$'s converge;

 foreach s do

$| \text{Policy evaluation}$

$| \pi(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a) [R(s, a, s') + \gamma V_\pi(s')];$

 | end

until $\pi(s)$'s converge;

Random initialize a policy
Let's say all goes down!



小吃 0	資電 0	排球 0
綜二 0	台達 0	籃球 0
總圖 0	工三 0	西門 0

After initialization of V_π

Input: MDP ($\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty$)

Output: $\pi(s)$'s for all s 's

For each state s , initialize $\pi(s)$ randomly;
repeat

 For each state s , initialize $V_\pi(s) \leftarrow 0$;

 repeat

 foreach s do

$V_\pi(s) \leftarrow \sum_{s'} P(s'|s; \pi(s)) [R(s, \pi(s), s') + \gamma V_\pi(s')]$;

 end

 until $V_\pi(s)$'s converge;

 foreach s do

 Policy improvement

$\pi(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a) [R(s, a, s') + \gamma V_\pi(s')]$;

 end

until $\pi(s)$'s converge;



Policy evaluation

Policy improvement



小吃 -∞	資電 -2	排球 -∞
綜二 -∞	台達 -1	籃球 -∞
總圖 -∞	工三 0	西門 -∞

After Policy Evaluation

Input: MDP ($\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty$)

Output: $\pi(s)$'s for all s 's

For each state s , initialize $\pi(s)$ randomly;
repeat

 For each state s , initialize $V_\pi(s) \leftarrow 0$;

 repeat

 foreach s do

$V_\pi(s) \leftarrow \sum_{s'} P(s'|s; \pi(s)) [R(s, \pi(s), s') + \gamma V_\pi(s')]$;

 end

 until $V_\pi(s)$'s converge;

 foreach s do

$\pi(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a) [R(s, a, s') + \gamma V_\pi(s')]$;

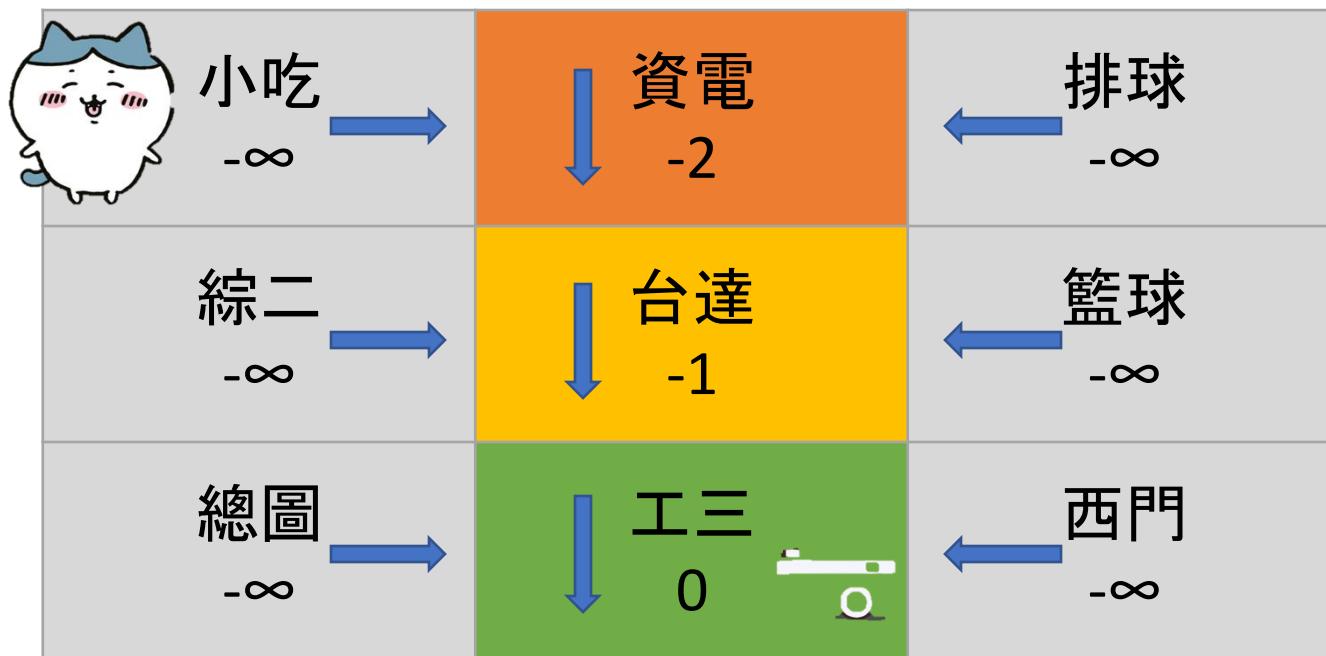
 end

until $\pi(s)$'s converge;

Policy evaluation

Policy improvement

After Policy Improvement



Input: MDP ($\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty$)

Output: $\pi(s)$'s for all s 's

For each state s , initialize $\pi(s)$ randomly;
repeat

 For each state s , initialize $V_\pi(s) \leftarrow 0$;

 repeat

 foreach s do

$V_\pi(s) \leftarrow \sum_{s'} P(s'|s; \pi(s)) [R(s, \pi(s), s') + \gamma V_\pi(s')]$;

 end

 until $V_\pi(s)$'s converge;

 foreach s do

 Policy improvement

$\pi(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a) [R(s, a, s') + \gamma V_\pi(s')]$;

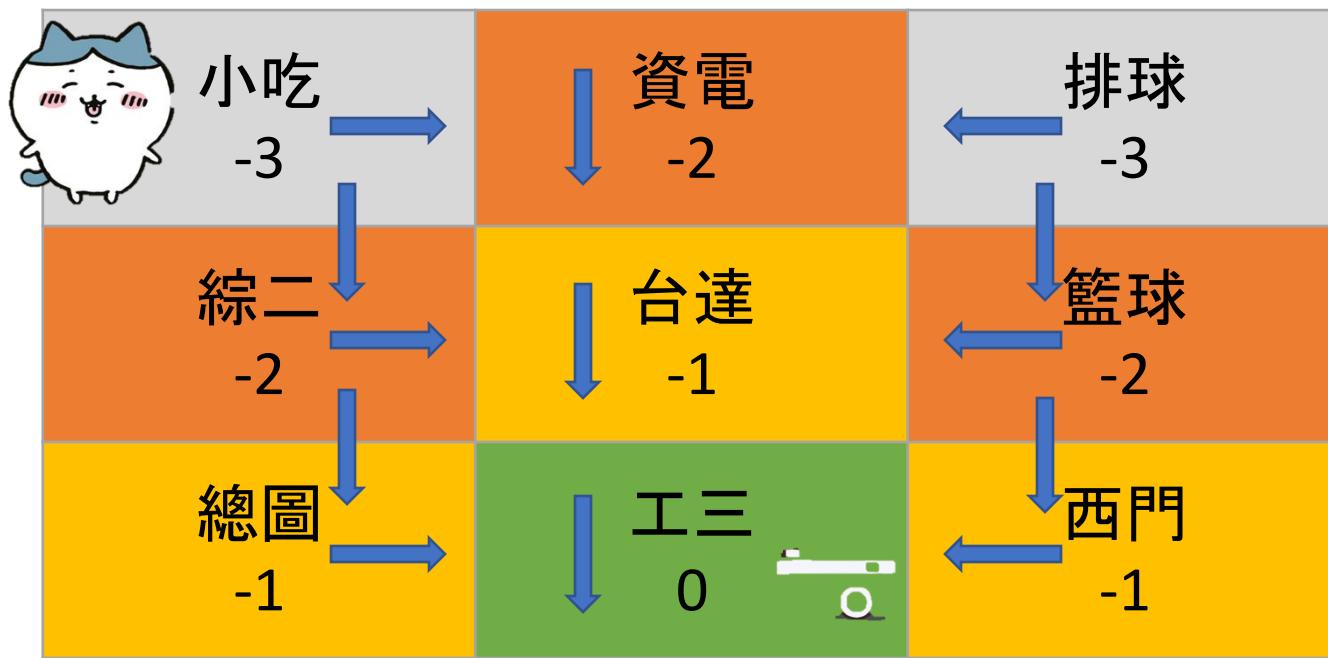
 end

until $\pi(s)$'s converge;

$$V(\text{籃球}) = \text{reward} + V(\text{西門}) = -1 + -\infty$$

$$\arg \max_a V(\text{籃球}) = \text{reward} + V(\text{台達}) = -1 + -1$$

$$V(\text{籃球}) = \text{reward} + V(\text{排球}) = -1 + -\infty$$



Input: MDP ($\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty$)
Output: $\pi(s)$'s for all s 's

For each state s , initialize $\pi(s)$ randomly;
repeat

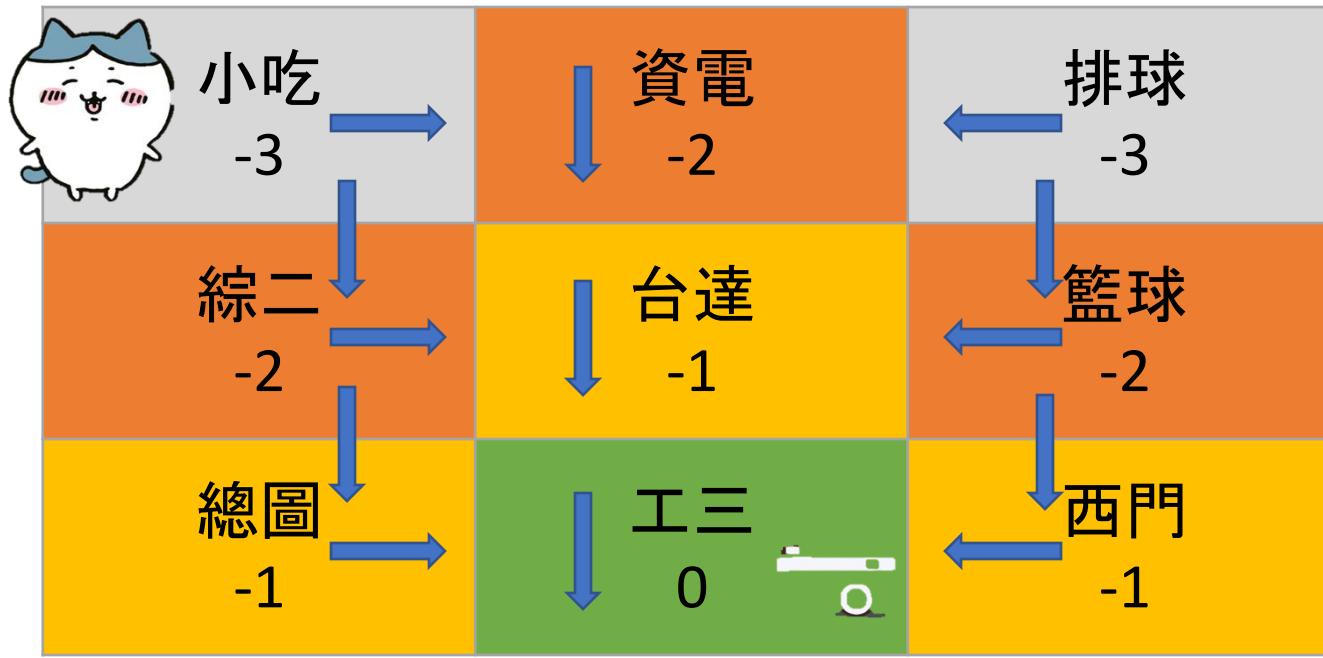
```

    For each state  $s$ , initialize  $V_\pi(s) \leftarrow 0$ ;
    repeat
        foreach  $s$  do
            |  $V_\pi(s) \leftarrow \sum_{s'} P(s'|s; \pi(s)) [R(s, \pi(s), s') + \gamma V_\pi(s')]$ ;
        end
    until  $V_\pi(s)$ 's converge;
    foreach  $s$  do
        |  $\pi(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a) [R(s, a, s') + \gamma V_\pi(s')]$ ;
    end
until  $\pi(s)$ 's converge;
```

Policy evaluation

Policy improvement

Policy Evaluation Again!



Input: MDP ($\mathbb{S}, \mathbb{A}, P, R, \gamma, H \rightarrow \infty$)

Output: $\pi(s)$'s for all s 's

For each state s , initialize $\pi(s)$ randomly;
repeat

For each state s , initialize $V_\pi(s) \leftarrow 0$;

repeat

foreach s do

Policy evaluation

$$V_\pi(s) \leftarrow \sum_{s'} P(s'|s; \pi(s)) [R(s, \pi(s), s') + \gamma V_\pi(s')];$$

end

until $V_\pi(s)$'s converge;

foreach s do

Policy improvement

$$\pi(s) \leftarrow \arg \max_a \sum_{s'} P(s'|s; a) [R(s, a, s') + \gamma V_\pi(s')];$$

end

until $\pi(s)$'s converge;

Policy Improvement.
Nothing Changed!
Converge!!



Did agent Interact with the Environment?

- No ! We model every transition and every reward
- But it is impossible to solve more complex problems like Flappy Bird
- We need model-free algorithms
 - Q-Learning
 - SARSA

Outline

- Homework
- Markov Decision Process (MDP)
 - Value Iteration
 - Policy Iteration
- Q-Learning & SARSA

Q-Learning

- Flappy bird



Q-Learning

- Flappy Bird
- States: $(\Delta x, \Delta y)$
- Actions: { fly, none }
- Reward:
 - +1: pass through a pipe
 - -5: die



Q-Learning

- Q-table(finite):

狀態	飛	不飛
$(\Delta x_1, \Delta y_1)$	1	20
$(\Delta x_1, \Delta y_2)$	20	-100
...
$(\Delta x_m, \Delta y_{n-1})$	-100	2
$(\Delta x_m, \Delta y_n)$	50	-200



- Update rule:
$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

Q-Learning

- Algorithm

Q-learning: An off-policy TD control algorithm

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

 Initialize S

 Repeat (for each step of episode):

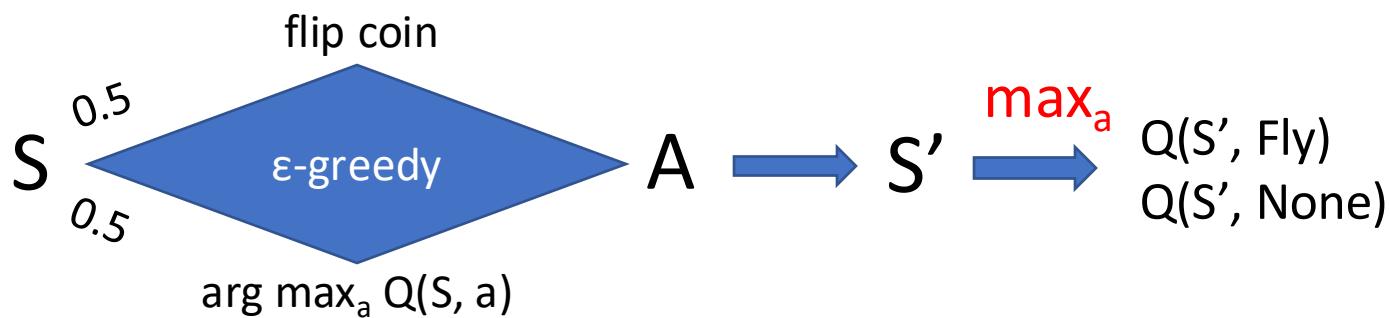
 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal



SARSA

- Algorithm

Sarsa: An on-policy TD control algorithm

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

 Initialize S

 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Repeat (for each step of episode):

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$$

$S \leftarrow S'; A \leftarrow A'$;

 until S is terminal

SARSA

- Q-table(finite):

狀態	飛	不飛
$(\Delta x_1, \Delta y_1)$	1	20
$(\Delta x_1, \Delta y_2)$	20	-100
...
$(\Delta x_m, \Delta y_{n-1})$	-100	2
$(\Delta x_m, \Delta y_n)$	50	-200



- Update rule:
$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$$

SARSA

- Algorithm

Sarsa: An on-policy TD control algorithm

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$
Repeat (for each episode):

 Initialize S

 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Repeat (for each step of episode):

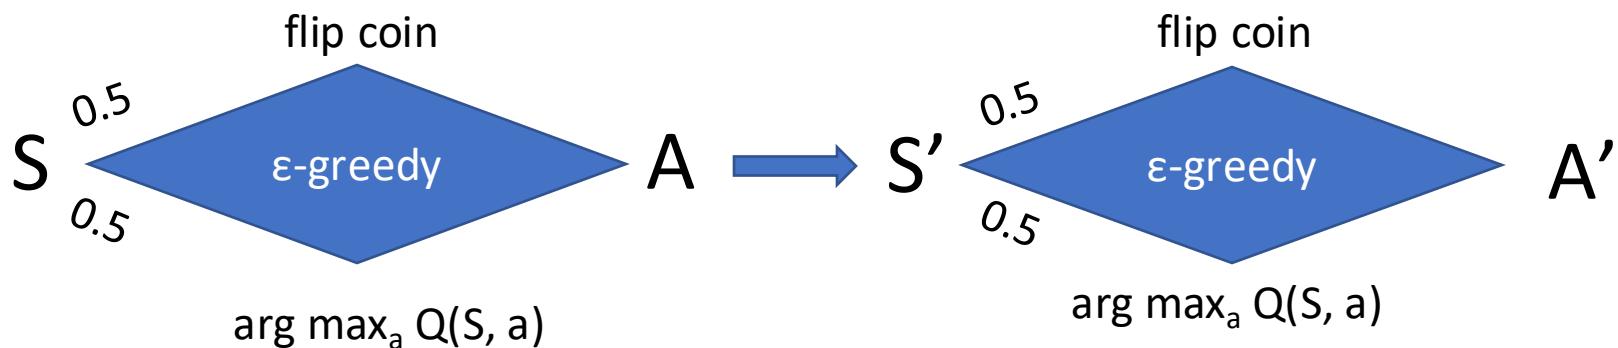
 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$$

$S \leftarrow S'; A \leftarrow A'$;

 until S is terminal



Q-Learning VS. SARSA

- Difference

```

Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$ 
Repeat (for each episode):
    Initialize  $S$  ②
    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
    Repeat (for each step of episode):
        Take action  $A$ , observe  $R, S'$ 
        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
         $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$ 
         $S \leftarrow S'; A \leftarrow A'$  ①
    until  $S$  is terminal
③

```

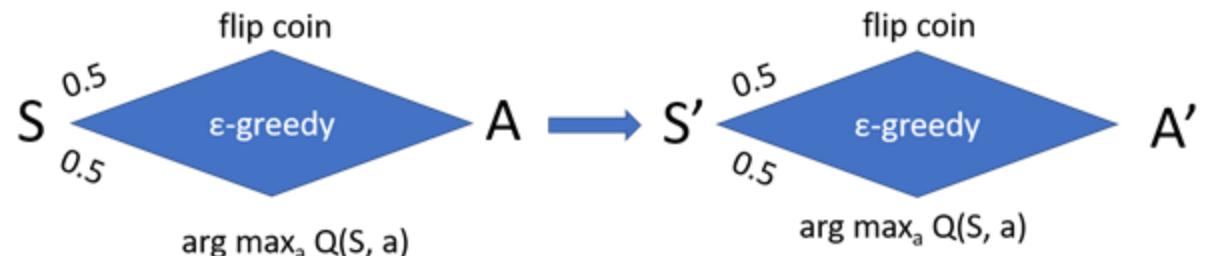


Figure 6.9: Sarsa: An on-policy TD control algorithm.

```

Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$ 
Repeat (for each episode):
    Initialize  $S$ 
    Repeat (for each step of episode):
        Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
        Take action  $A$ , observe  $R, S'$ 
         $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$ 
         $S \leftarrow S';$  what about  $A'$ ? ①
    until  $S$  is terminal
③
②
① this is like following a greedy policy (e.g.  $\epsilon=0$ , NO exploration)

```

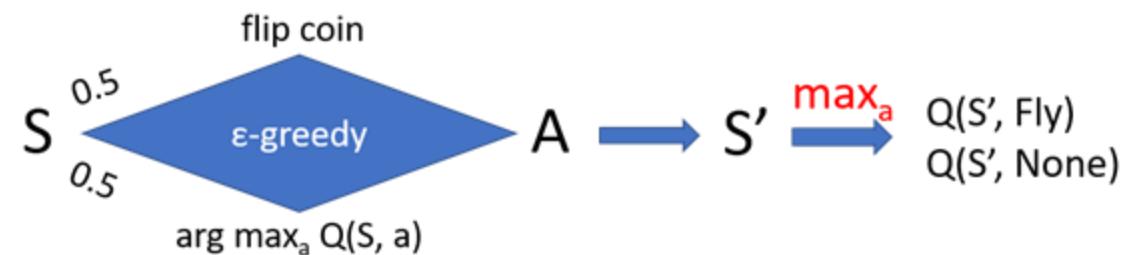
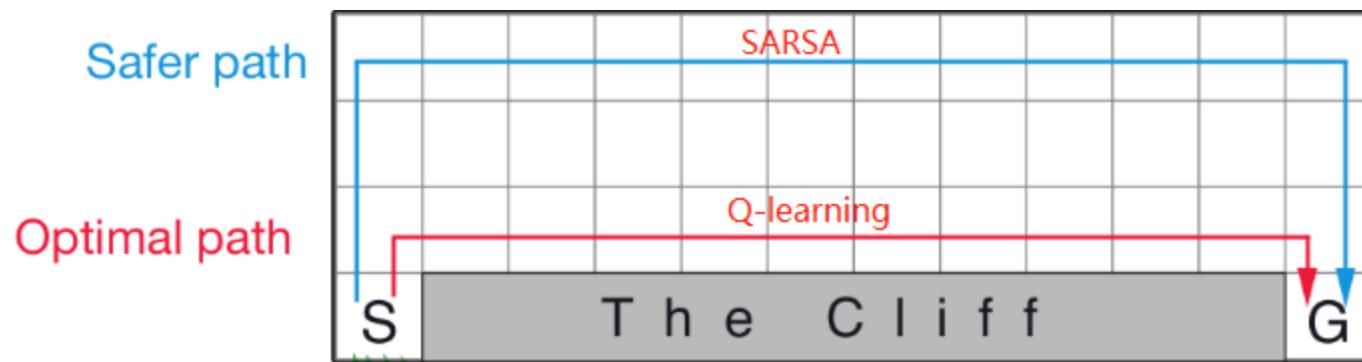


Figure 6.12: Q-learning: An off-policy TD control algorithm.

Q-Learning VS. SARSA

- Cliff Walking



Thanks! Be a Happy Bird